



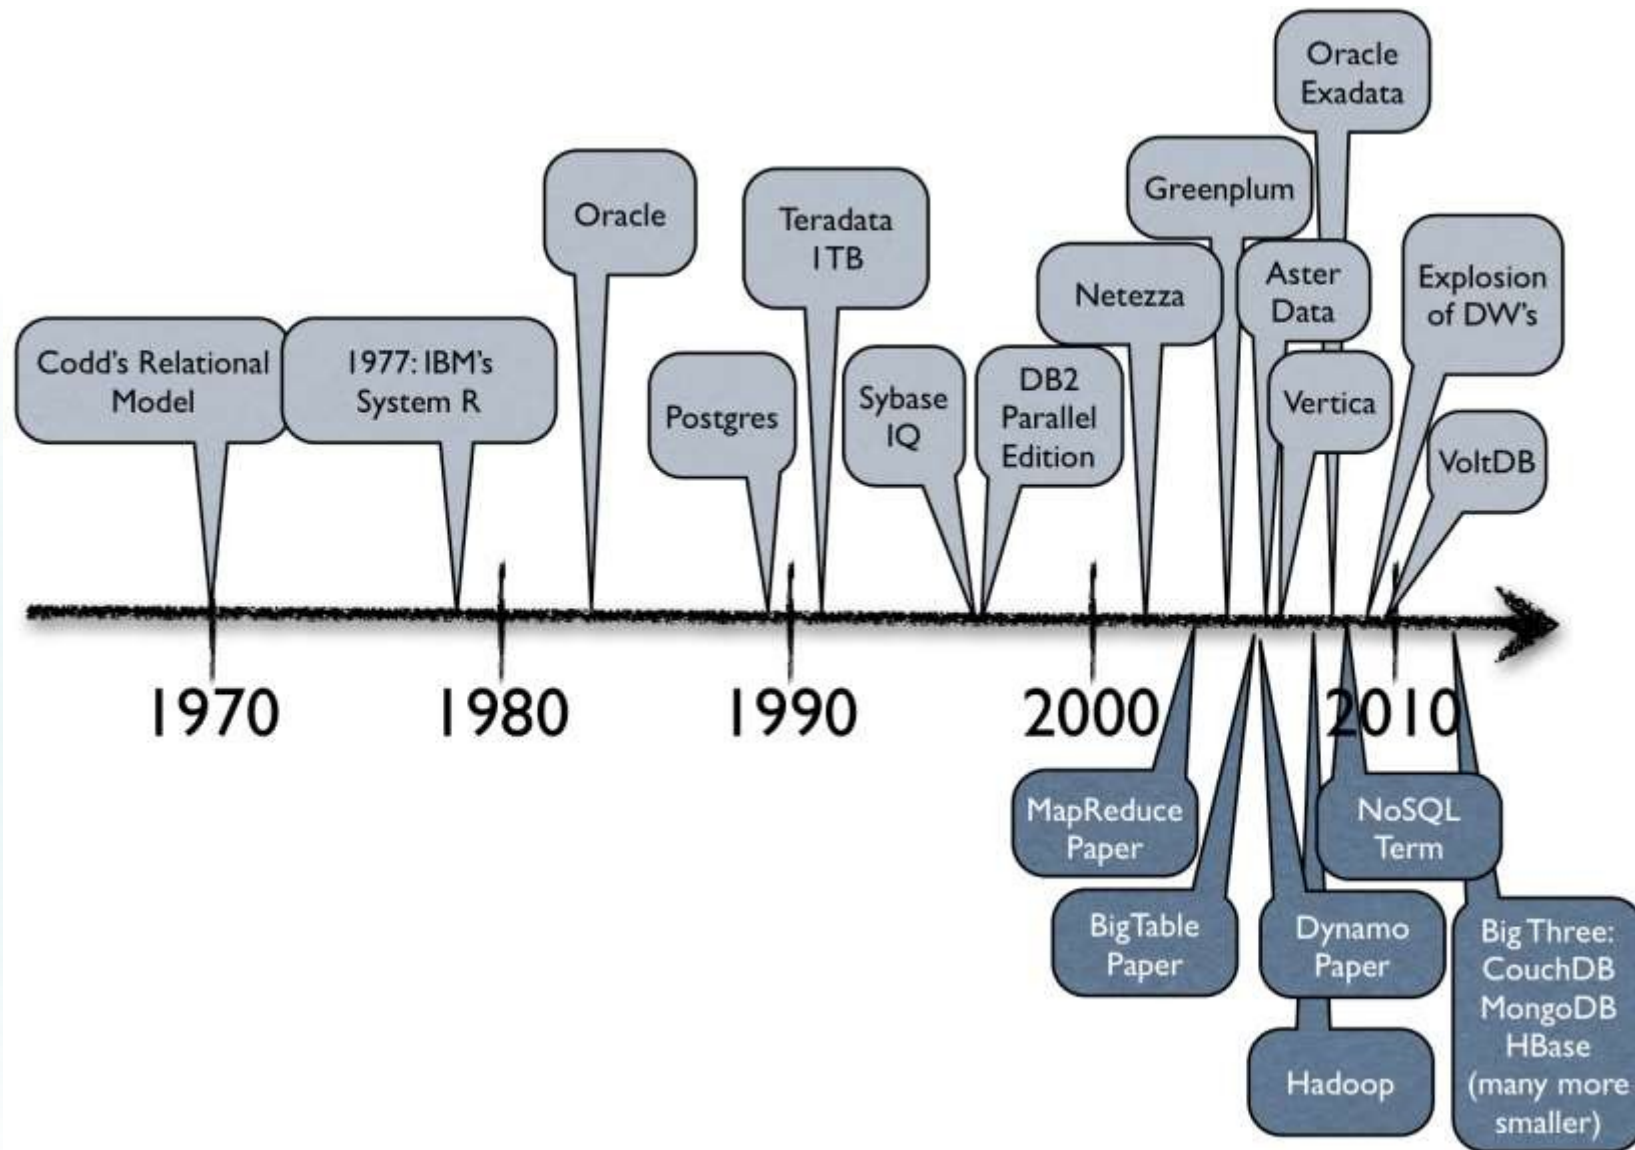
An introduction to **NoSQL** databases

Dr. Faisal Kamiran

Agenda:

- ▶ A brief history of databases
- ▶ NoSQL why, what and when?
- ▶ Aggregate Data Models
- ▶ BASE vs ACID
- ▶ CAP theorem
- ▶ Denormalization

A brief history of databases



Relational databases

Benefits of Relational databases:

- Designed for all purposes
- ACID
- Strong consistency, concurrency, recovery
- Mathematical background
- Standard Query language (SQL)
- Lots of tools to use with i.e: Reporting services, entity frameworks, ...
- Vertical scaling (up scaling)


Era of Distributed Computing

But...

- ❑ Relational databases were not built for **distributed applications.**

Because...

- ❑ Joins are expensive
- ❑ Hard to scale horizontally
- ❑ Expensive (product cost, hardware, Maintenance)



Era of Distributed Computing

But...

- ☐ Relational databases were not built for **distributed applications.**


Because...

- ☐ Joins are expensive
- ☒ **Hard to scale horizontally**
- ☐ Expensive (product cost, hardware, Maintenance)

And....

It's weak in:

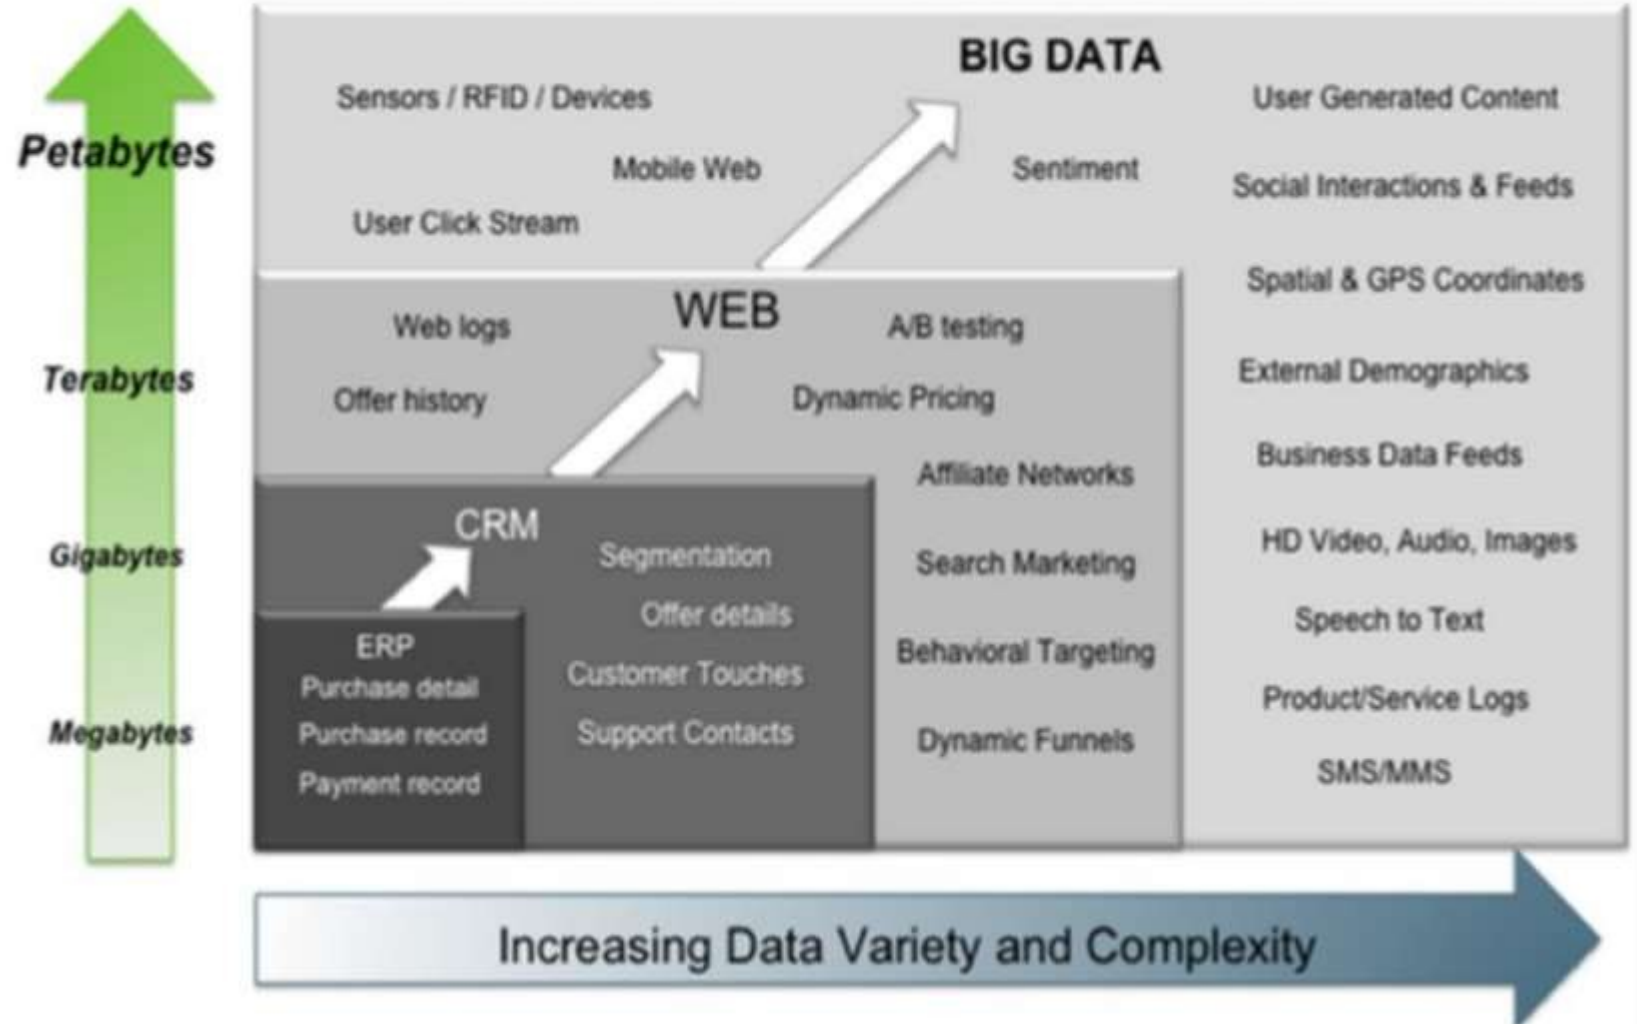
- ☐ Speed (performance)
- ☐ High availability
- ☐ Partition tolerance



Rise of Big data

Three V(s) of Bigdata:

- ▶ Volume
- ▶ Velocity
- ▶ Variety



NoSQL why, what and when?

- ▶ Google & Amazon built their own databases (Big table & Dynamo)
- ▶ Facebook invented Cassandra and is using thousands of them
- ▶ #NoSQL was a twitter hashtag for a conference in 2009
- ▶ The name doesn't indicate its characteristics
- ▶ There is no strict definition for NoSQL databases
- ▶ There are more than 150 NoSQL databases (nosql-database.org)



Characteristics of NoSQL databases

- ▶ Non relational
- ▶ Cluster friendly
- ▶ Schema-less
- ▶ 21 century web
- ▶ Open-source



Characteristics of NoSQL databases

NoSQL avoids:

- ▶ Overhead of ACID transactions
- ▶ Complexity of SQL query
- ▶ Burden of up-front schema design
- ▶ DBA presence
- ▶ Transactions (It should be handled at application layer)

Provides:

- ▶ Easy and frequent changes to DB
- ▶ Horizontal scaling (scaling out)
- ▶ Solution to Impedance mismatch
- ▶ Fast development



NoSQL is getting more & more popular

The Google logo, featuring the word "Google" in its signature multi-colored font.The Facebook logo, consisting of the word "facebook" in a blue, lowercase, sans-serif font.The Amazon logo, featuring the word "amazon" in a black, lowercase, sans-serif font with a curved orange arrow underneath.The LinkedIn logo, featuring the word "LinkedIn" in a black, sans-serif font, with the "in" part enclosed in a blue square.The Yahoo! logo, featuring the word "YAHOO!" in a purple, serif font.The Twitter logo, featuring a blue bird icon above the word "twitter" in a blue, lowercase, sans-serif font.The Netflix logo, featuring the word "NETFLIX" in a white, bold, sans-serif font on a red rectangular background.The eBay logo, featuring the word "ebay" in a multi-colored, lowercase, sans-serif font.The Guardian logo, featuring the word "theguardian" in a blue, lowercase, sans-serif font, followed by a red and blue donkey icon with white stars.

What is a schema-less data model?

In relational Databases:

- ▶ You can't add a record which does not fit the schema
- ▶ You need to add NULLs to unused items in a row
- ▶ We should consider the datatypes. i.e : you can't add a string to an integer field
- ▶ You can't add multiple items in a field (You should create another table: primary-key, foreign key, joins, normalization, ... !!!)

```
create table customers (id int, firstname text, lastname text)  
insert into customers (firstname, middlename, lastname) values (...)
```




What is a schema-less datamodel?

In NoSQL Databases:

- ▶ There is no schema to consider
- ▶ There is no unused cell
- ▶ There is no datatype (implicit)
- ▶ Most of considerations are done in application layer
- ▶ We gather all items in an aggregate (document)

Relational Model



Document Model

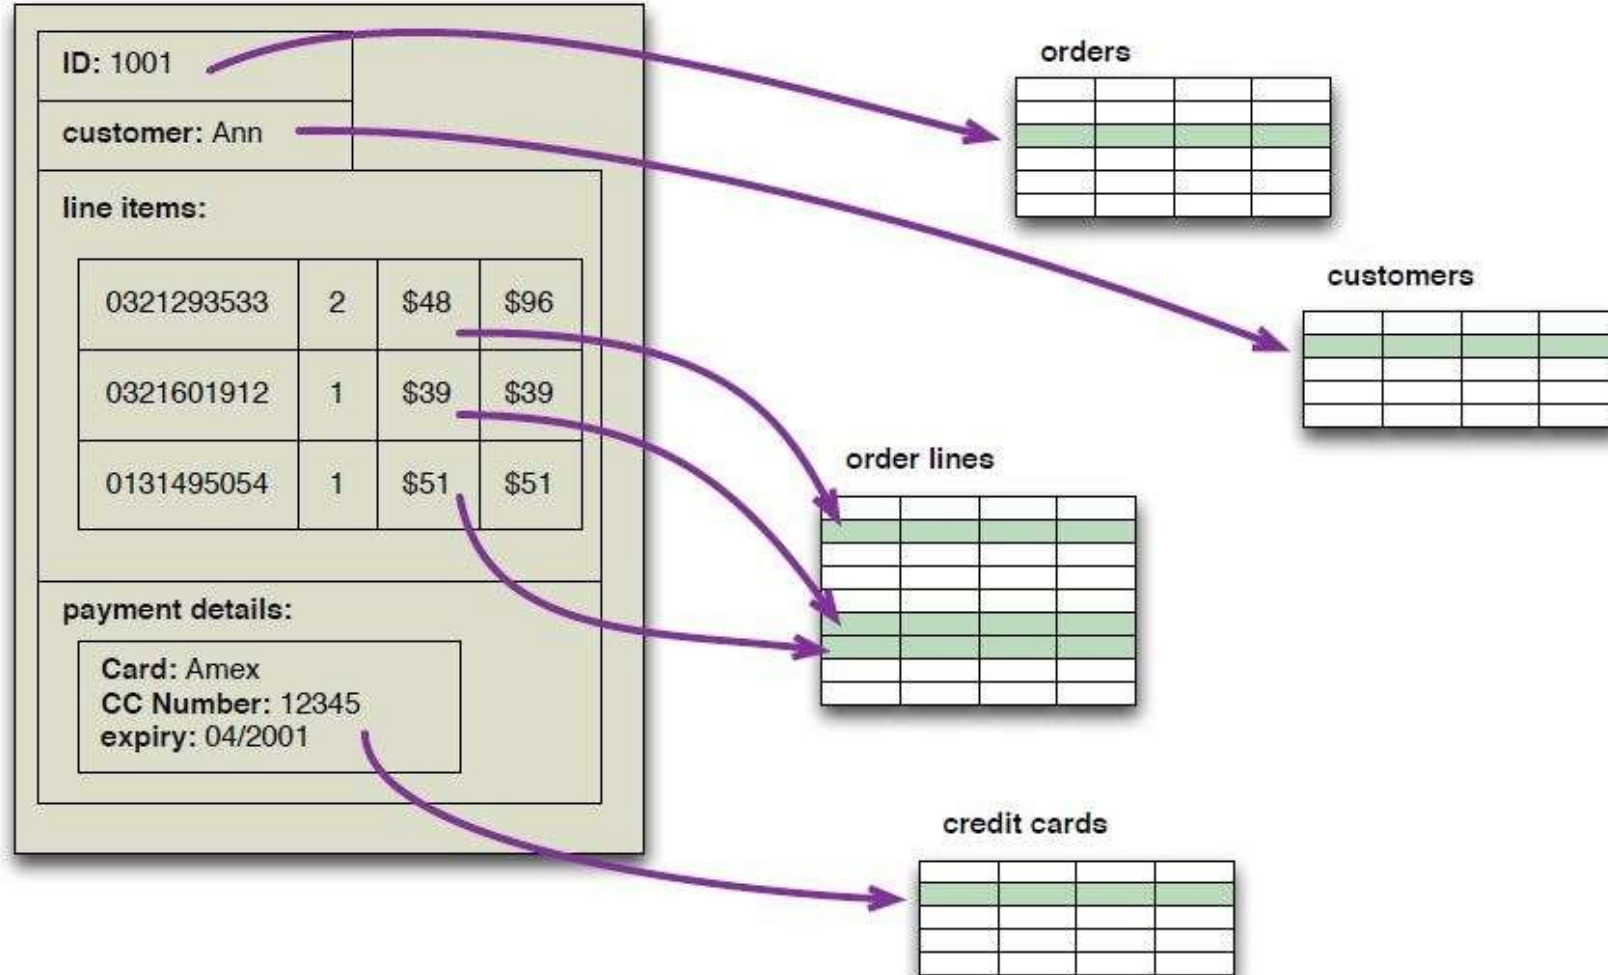
Collection ("Things")



What is Aggregation?

- ▶ The term comes from Domain Driven Design
- ▶ Shared nothing architecture
- ▶ An aggregate is a cluster of domain objects that can be treated as a single unit
- ▶ Aggregates are the basic element of transfer of data storage - you request to load or save whole aggregates
- ▶ Transactions should not cross aggregate boundaries
- ▶ This mechanism reduces the join operations to a minimal level

What is Aggregation?



What is Aggregation?

```
{
  "id": "1001",
  "firstName": "Ann",
  "lastName": "Williams",
  "age": 55,
  "purchasedItems":
  {
    0321290533 {qty, price...}
    0321601912 {qty, price...}
    0131495054 {qty, price...}
  }
  "paymentDetails":
  { cc info... }
  "address":
  {
    "street": "1234 Park",
    "city": "San Francisco",
    "state": "CA",
    "zip": "94102"
  }
}
```

ID: 1001			
customer: Ann			
line items:			
0321293533	2	\$48	\$96
0321601912	1	\$39	\$39
0131495054	1	\$51	\$51
payment details:			
Card: Amex CC Number: 12345 expiry: 04/2001			

aggregate



What is Aggregation?



Aggregate Data Models

NoSQL databases are classified in four major datamodels:

- ▶ Key-value
- ▶ Document
- ▶ Column family
- ▶ Graph

Each DB has its own query language



Key-value data model

- ▶ The main idea is the use of a hash table
- ▶ Access data (values) by strings called keys
- ▶ Data has no required format – data may have any format
- ▶ Data model: (key, value) pairs
- ▶ Basic Operations:
Insert(key,value), Fetch(key),Update(key), Delete(key)

Car	
Key	Attributes
1	Make: Nissan Model: Pathfinder Color: Green Year: 2003
2	Make: Nissan Model: Pathfinder Color: Blue Color: Green Year: 2005 Transmission: Auto

Key-value data model

- ▶ “Value” is stored as a “blob”
 - Without caring or knowing what is inside
 - Application is responsible for understanding the data
- ▶ Main observation from Amazon (using **Dynamo**)
 - “There are many services on Amazon’s platform that only need primary-key access to a data store.”

E.g. Best seller lists, shopping carts, customer preferences, session management, sales rank, product catalog



Column family data model

- ▶ The column is lowest/smallest instance of data.
- ▶ It is a tuple that contains a name, a value and a timestamp

ColumnFamily: Authors		
Key	Value	
"Eric Long"	Columns	
	Name	Value
	"email"	"eric (at) long.com"
	"country"	"United Kingdom"
	"registeredSince"	"01/01/2002"
"John Steward"	Columns	
	Name	Value
	"email"	"john.steward (at) somedomain.com"
	"country"	"Australia"
	"registeredSince"	"01/01/2009"
"Ronald Mathies"	Columns	
	Name	Value
	"email"	"ronald (at) sodeso.nl"
	"country"	"Netherlands, The"
	"registeredSince"	"01/01/2010"

Column family data model

Some statistics about Facebook Search (using **Cassandra**)

- ❖ MySQL > 50 GB Data
 - Writes Average : ~300 ms
 - Reads Average : ~350 ms
- ❖ Rewritten with Cassandra > 50 GB Data
 - Writes Average : 0.12 ms
 - Reads Average : 15 ms



Document-based datamodel

- ▶ Usually JSON like interchange model.
- ▶ Query Model: JavaScript-like or custom.
- ▶ Aggregations: **Map/Reduce**
- ▶ Indexes are done via B-Trees.
- ▶ unlike simple key-value stores, both keys and values are fully searchable in document databases.

```
{
  person: {
    first_name: "Peter",
    last_name: "Peterson",
    addresses: [
      {street: "123 Peter St"},
      {street: "504 Not Peter St"}
    ]
  }
}
```



Overview of a Document-based datamodel

```
SELECT name, pic, profile_url
FROM user
WHERE uid = me()
```



```
SELECT message, attachment
FROM stream
WHERE source_id = me() AND type = 80
```

```
SELECT name
FROM friendlist
WHERE owner = me()
```



```
SELECT name
FROM group
WHERE gid IN ( SELECT gid
                FROM group_member
                WHERE uid = me() )
```



```
SELECT name, pic
FROM user
WHERE online_presence = "active"
AND
uid IN ( SELECT uid2
          FROM friend
          WHERE uid1 = me() )
```

A sample MongoDB query

MySQL

```
SELECT * FROM customerOrder, orderItem, product  
WHERE  
customerOrder.orderId = orderItem.customerOrderId  
AND orderItem.productId = product.productId  
AND product.name LIKE '%Refactoring%'
```

MongoDB

```
db.orders.find( {"items.product.name":/Refactoring/})
```

**There is no join in MongoDB query
Because we are using an aggregate data model**

BASE

Almost the opposite of ACID.

- ▶ Basically available: Nodes in the a distributed environment can go down, but the whole system shouldn't be affected.
- ▶ Soft State (scalable): The state of the system and data changes over time.
- ▶ Eventual Consistency: Given enough time, data will be consistent across the distributed system.

BASE vs ACID

ACID:

- Strong consistency.
- Less availability.
- Pessimistic concurrency.
- Complex.

BASE:

- Availability is the most important thing. Willing to sacrifice for this (CAP).
- Weaker consistency (Eventual).
- Best effort.
- Simple and fast.
- Optimistic.

What we need?

- ▶ We need a distributed database system having such features:
 - **Fault tolerance**
 - ▶ – **High availability**
 - ▶ – **Consistency**
 - ▶ – **Scalability**

What we need?

- ▶ We need a distributed database system having such features:
 - **Fault tolerance**
 - **High availability**
 - **Consistency**
 - **Scalability**

Which is impossible!!!

According to **CAP theorem**

Should we...?

- ☐ In some cases getting an answer quickly is more important than getting a correct answer
- ☐ By giving up ACID properties, one can achieve higher performance and scalability.
- ☐ Any data store can achieve Atomicity, Isolation and Durability but do you always need **consistency**?
- ☐ Maybe we should implement Asynchronous Inserts and updates and should not wait for confirmation?

CAP theorem

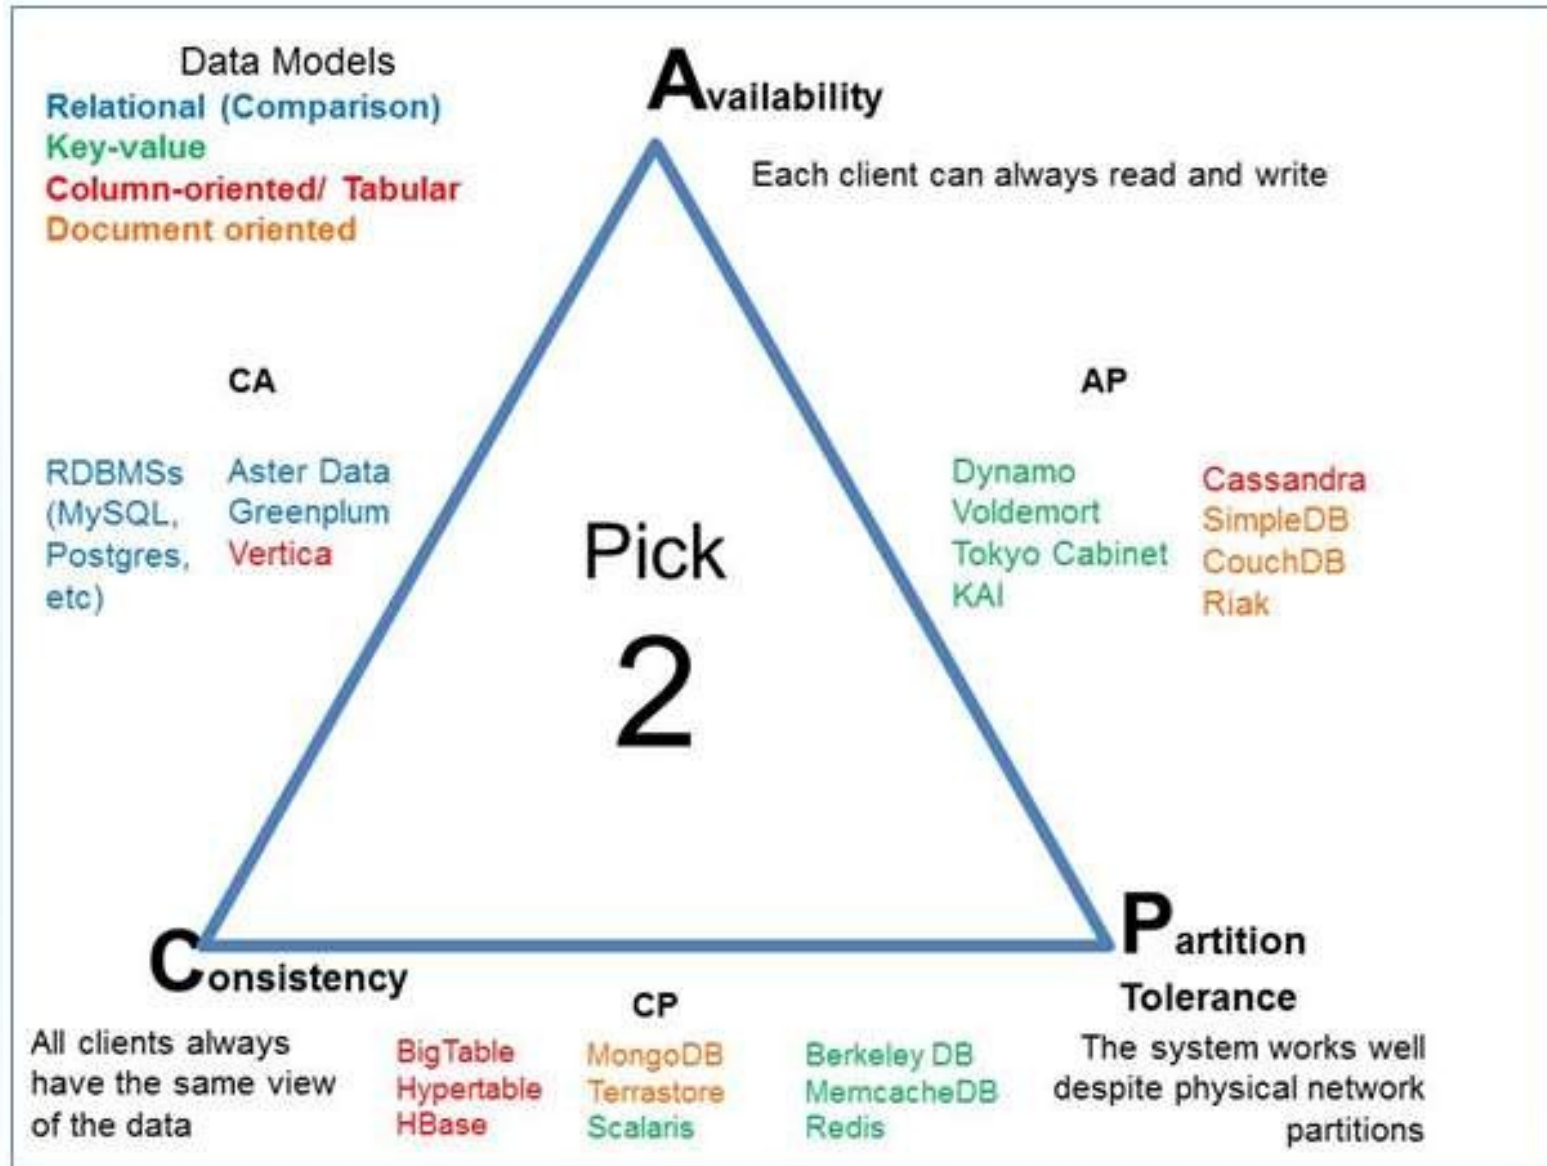
- ❑ **Consistency**: **C**lients should read the same **d**ata. There are many levels of consistency.
- ❑ **Availability**: Data to be available.
- ❑ **Partial Tolerance**: Data to be partitioned across network segments due to network failures.

2000 Prof. Eric Brewer, PoDC Conference Keynote

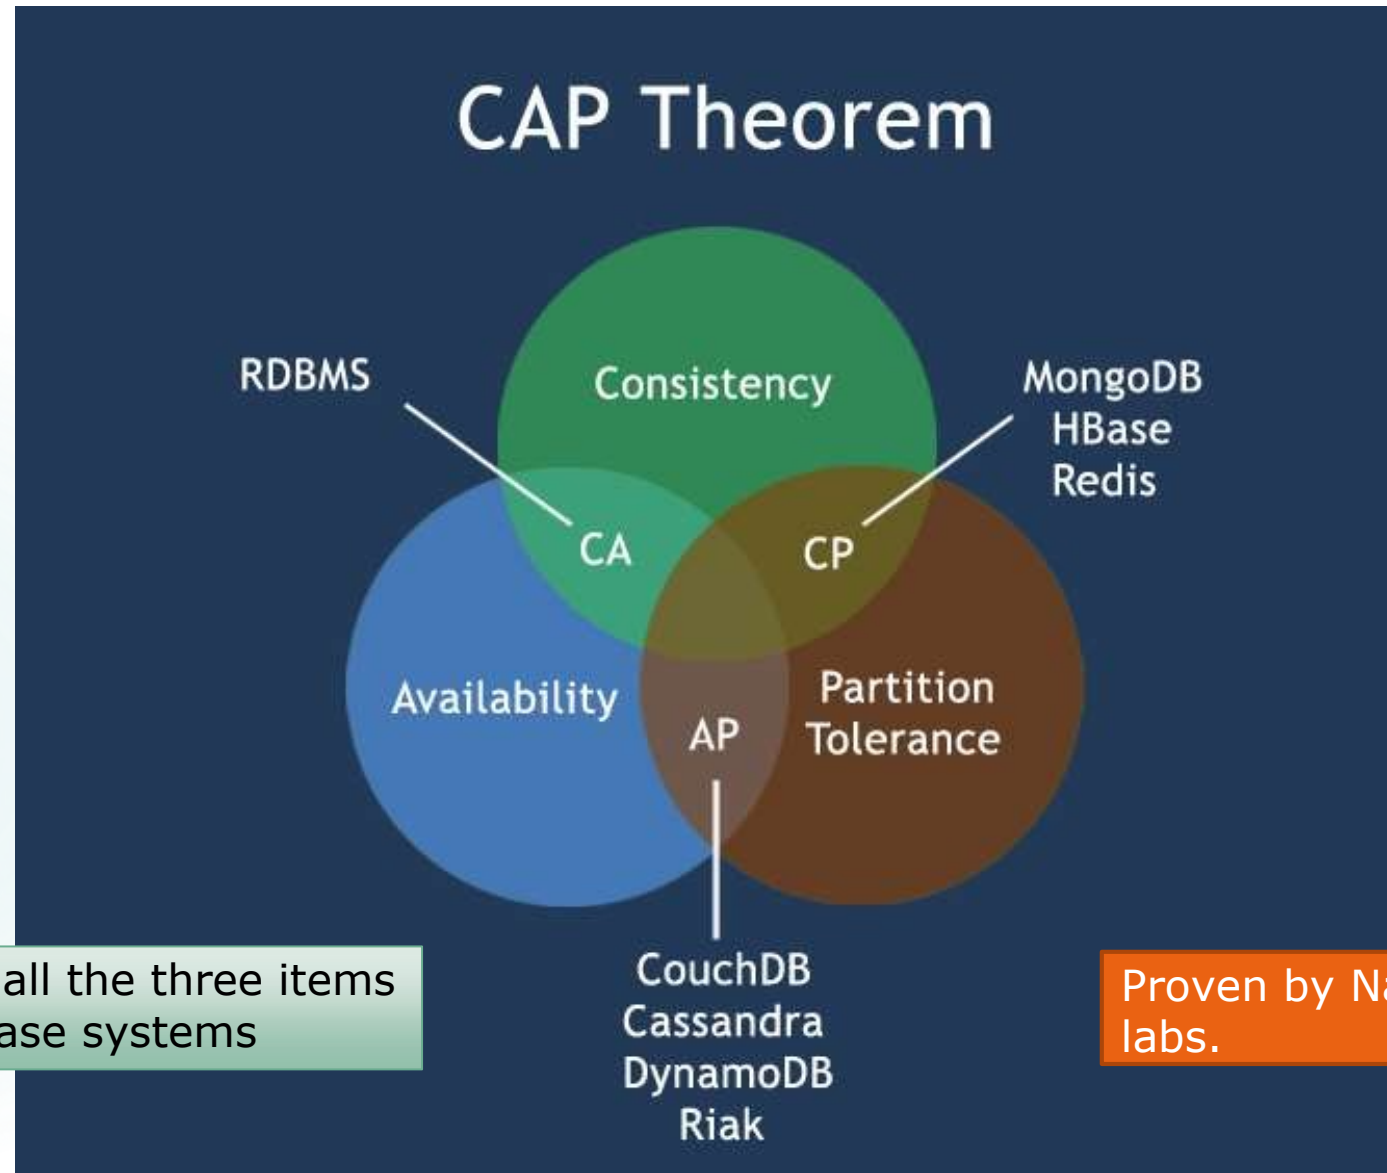
2002 Seth Gilbert and Nancy Lynch, ACM SIGACT News 33(2)

“ Of three properties of shared-data systems - data **C**onsistency, system **A**vailability and tolerance to network **P**artitions - only two can be achieved at any given moment in time.”

Which data model to choose



CAP theorem in different SQL/NoSQL databases




We can not achieve all the three items
In distributed database systems
(center)

Proven by Nancy Lynch et al. MIT
labs.


CAP theorem : A simple proof

Consistent and available
No partition.




CAP theorem : A simple proof

Available and partitioned
Not consistent, we get back old data.



CAP theorem : A simple proof

Consistent and partitioned
Not available, waiting...



WHEN TO USE A NOSQL DATABASE?

- Large amounts of data
 - Terabytes and Petabytes of data
- Need horizontal scalability
- Need high throughput – fast reads
- Need a flexible schema
 - No fixed number of columns
- Need high availability
- Need to be able to store different data type formats
- Users are distributed – low latency
- Need redundancy in case of failures

WHEN NOT TO USE A NOSQL DATABASE?

- Need ACID Transactions
- Need ability to do JOINS
- Ability to do aggregations and analytics
- Have changing business requirements
- Queries are not available and need to have flexibility
- Have a small dataset

Denormalization

Definition :-

- ❖ Denormalization is the process of taking a normalized database and modifying table structures to allow controlled redundancy for increased database performance.

Method of denormalization:-


- 1) Adding Redundant Columns.
- 2) Adding Derived Columns
- 3) Combining Tables
- 4) Repeating Groups
- 5) Partitioning Relations

Adding Redundant Columns


You can add redundant columns to eliminate frequent joins.

For example, if frequent joins are performed on the *titleauthor* and *authors* tables in order to retrieve the author's last name, you can add the *au_lname* column to *titleauthor*.

```
select ta.title_id, a.au_id, a.au_lname  
from titleauthor ta, authors a  
where ta.au_id = a.au_id
```



```
select title_id, au_id, au_lname  
from titleauthor
```




- Adding redundant columns eliminates joins for many queries. The problems with this solution are that it:
- Requires maintenance of new column. All changes must be made to two tables, and possibly to many rows in one of the tables.
- Requires more disk space, since *au_lname* is duplicated.

Adding Derived Columns


Adding derived columns can help eliminate joins and reduce the time needed to produce aggregate values.

For example. Frequent joins are needed between the *titleauthor* and *titles* tables to provide the total advance for a particular book title.

```
select title, sum(advance)  
from titleauthor ta, titles t  
where ta.title_id = t.title_id  
group by title_id
```



```
select title, sum_adv  
from titles
```



- You can create and maintain a derived data column in the *titles* table, eliminating both the join and the aggregate at run time. This increases storage needs, and requires maintenance of the derived column whenever changes are made to the *titles* table.

Combining Tables


If most users need to see the full set of joined data from two tables, collapsing the two tables into one can improve performance by eliminating the join.

For example, users frequently need to see the author name, author ID, and the *blurbs* copy data at the same time. The solution is to collapse the two tables into one. The data from the two tables must be in a one-to-one relationship to collapse tables.

**select a.au_id, a.au_lname,
b.copy
from authors a, blurbs b
where a.au_id = b.au_id**

au_id	au_lname

au_id	copy

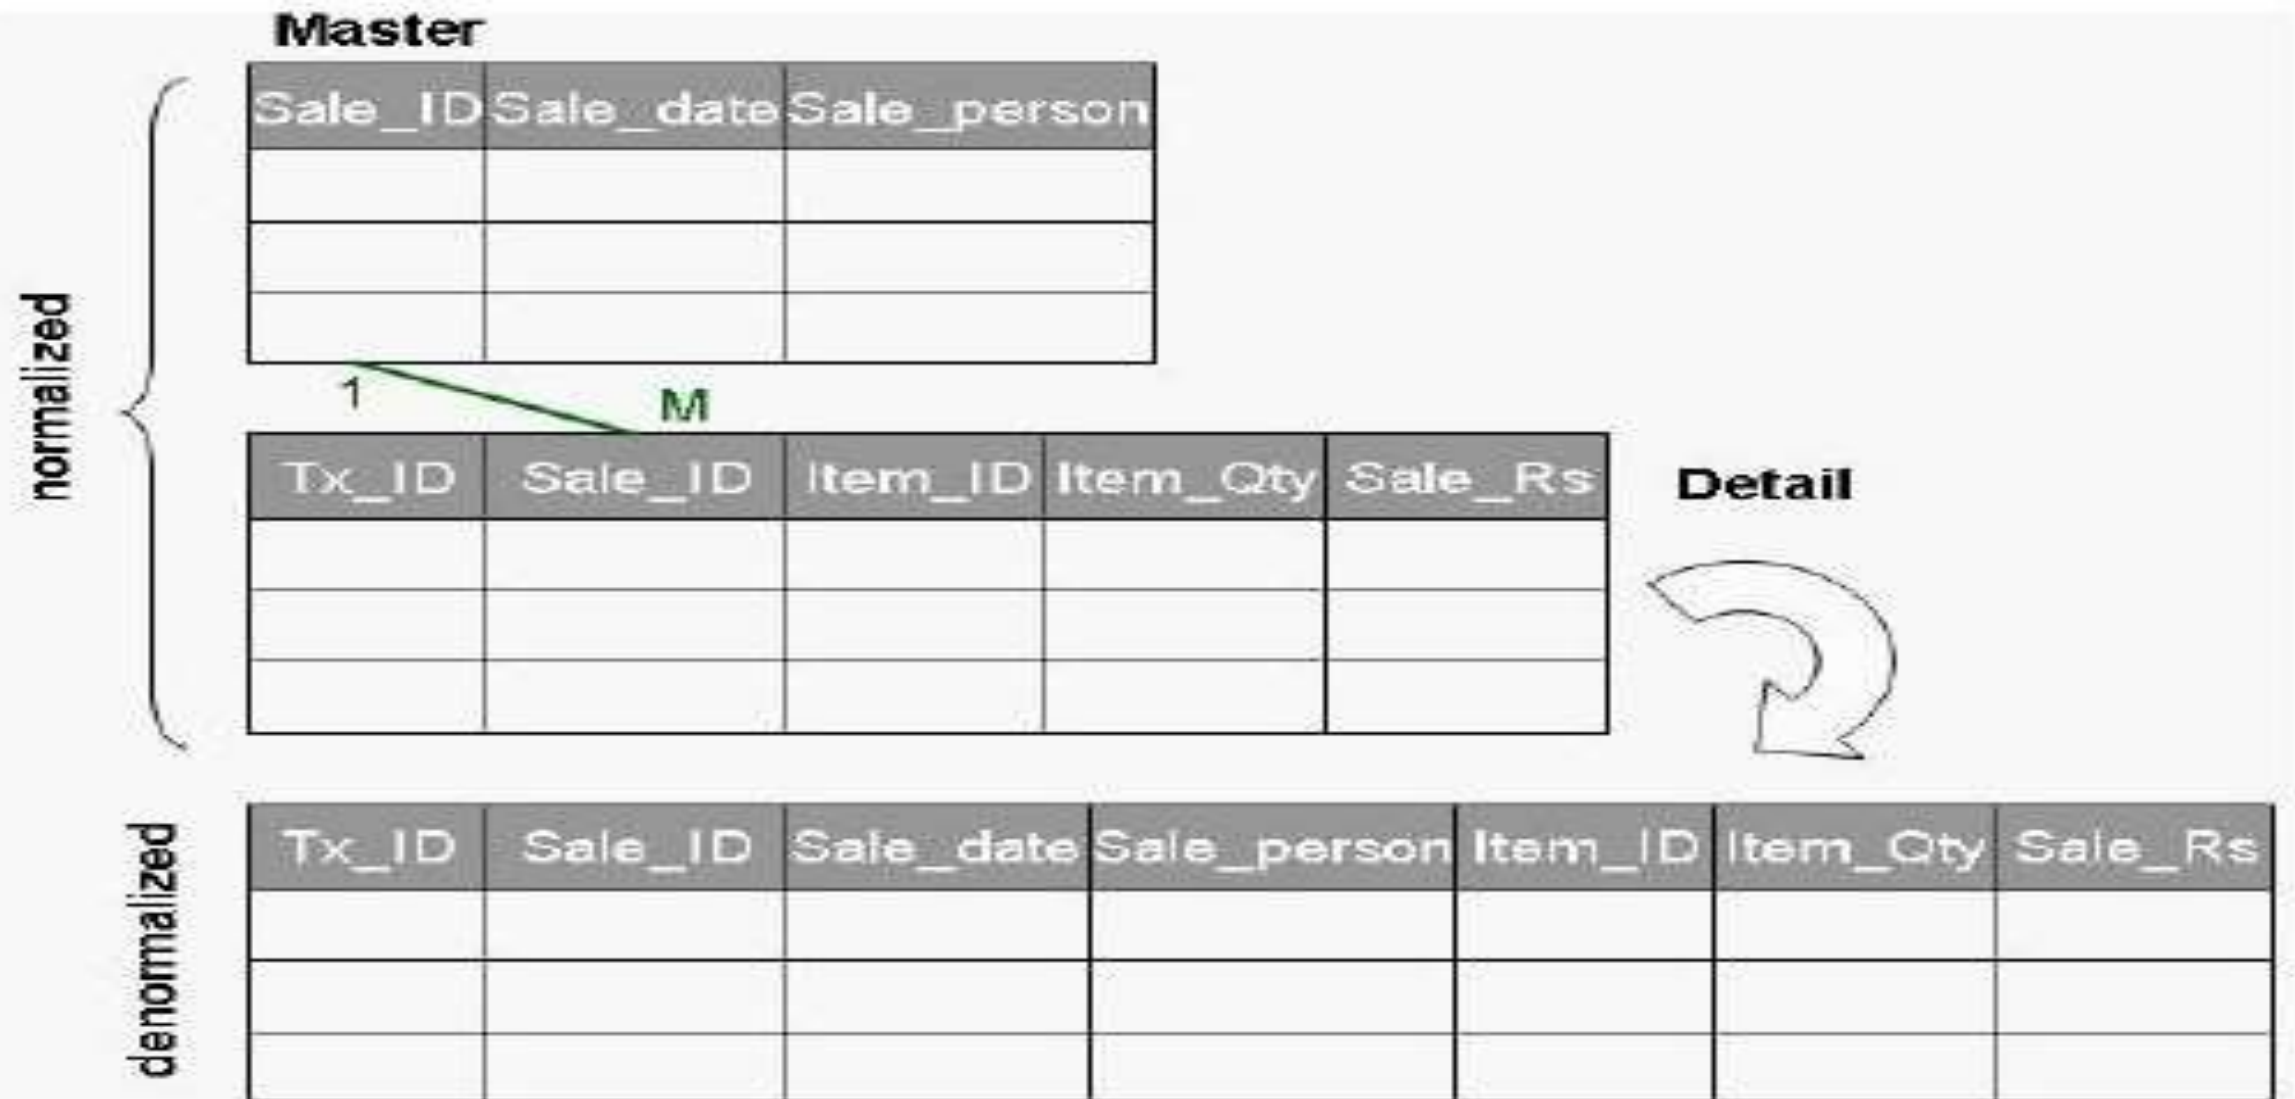


select * from newauthors

newauthors

au_id	au_lname	copy

More examples:-



Q & A



References

- Database design & Relational theory, C.J.Date
- Database system .8th edition.
- Data Normalization, Denormalization, and the Forces of Darkness a white paper by Melissa Hollingsworth.
- NoSQL distilled, Martin Fowler
- Martin Fowler's presentation at Goto conference
- www.mongodb.org
- Pooyan Mehrparvar